

Time Series Project

Liam Flaherty, Zhao Qu, and Zhijiang Yang

Professor Martin

NCSU: ST534-001

October 14, 2024

Contents

1	Introduction	3
2	Model Selection	4
2.1	White Noise Test	4
2.2	(Weak) Stationarity Test	4
2.3	Autocorrelation And Partial Autocorrelation Functions	5
2.4	Seasonality	6
2.5	Suggested Models	7
2.6	Model Diagnostics	8
3	Parameter Selection	9
3.1	Heart Rate Data	9
3.2	Weight Data	10
4	Forecasting	11
5	Appendix	13

1 Introduction

In 2022, I had a small surgery on my back that kept me pretty immobile for months on end. Once I was able to start moving normally again, my strength and endurance were both fractions of what they were prior to undergoing the procedure.

To build back my health, I started running and lifting weights. Progress could be judged by direct measurement (e.g. how far I could run or how much weight I could lift), but hopefully, changes in these direct measurements would also manifest themselves in proxy measurements (e.g. an increase in strength would lead to an increase in muscle mass and thus body weight, while an increase in endurance would lead to better cardiovascular health).

To that end, I made it a goal to increase my body weight at a constraint of a steady resting heart rate. With a few exceptions, I recorded progress in these areas each morning from August 2022 to March 2023 by using a blood pressure and heart rate gauge I bought from CVS and a bathroom scale. We reserve the data post February 2023 as the test set, and plot the training set in Figures 1.1-1.2 below.




Figure 1.1: Evolution Of Heart Rate




Figure 1.2: Evolution Of Weight

2 Model Selection

2.1 White Noise Test

It is clear from inspection that the time series for weight is not white noise. We use the Ljung-Box Q Test to determine whether we need to fit a model for the heart rate data. The function `Box.test()` in R provides a nice way to look at this. Upon running this function on our data, we get the below output in Figure 2.1. With a p-value of about 0.004, we reject the null hypothesis of “white noise” under a significance level of $\alpha = 0.05$.

```
> whitenoise6=Box.test(ts_hr,          #Do we need to fit model?#
+                        lag=6,
+                        type="Ljung-Box")
> whitenoise6                          #p small \implies yes#

Box-Ljung test

data:  ts_hr
X-squared = 19.069, df = 6, p-value = 0.004048
```

Figure 2.1: R Output For Ljung-Box Q Test

2.2 (Weak) Stationarity Test

We can test whether or not we need to take a difference in either dataset in order to make the data stationary (test if there is a unit root) with the Augmented Dickey-Fuller Test. At least visually, it appears there is trending in the weight dataset, but constant mean in the heart rate dataset. This is borne out by a formal test of the data, which we show in Figures 2.2-2.3 below. Note that there is not enough evidence to reject the null hypothesis of “non-stationary” for the weight data under a significance level of $\alpha = 0.05$.

```
> adf_result_weight=suppresswarnings(adf.test(ts_weight))
> adf_result_weight                      #difference needed#

Augmented Dickey-Fuller Test

data:  ts_weight
Dickey-Fuller = -2.8013, Lag order = 5, p-value = 0.2418
alternative hypothesis: stationary
```

Figure 2.2: ADF Output For Weight

```
> adf_result_hr=suppresswarnings(adf.test(ts_hr))
> adf_result_hr

Augmented Dickey-Fuller Test

data:  ts_hr
Dickey-Fuller = -4.775, Lag order = 5, p-value = 0.01
alternative hypothesis: stationary
```

Figure 2.3: ADF Output For Heart Rate

Indeed, after taking a difference (Figure 2.4), we see the time series appears significantly more stationary than previously (Figure 1.2).

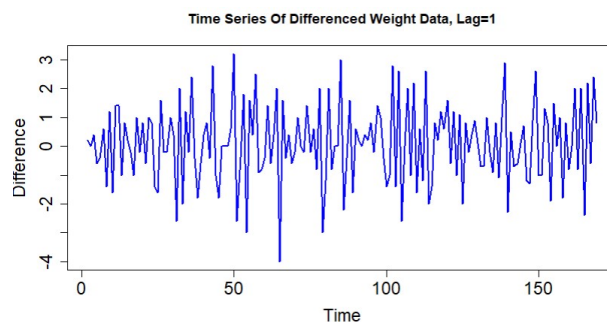


Figure 2.4: Time Series After Difference

2.3 Autocorrelation And Partial Autocorrelation Functions

The autocorrelation and partial autocorrelation for our heart rate data (Figure 2.5) and weight data (Figure 2.6) are shown below.

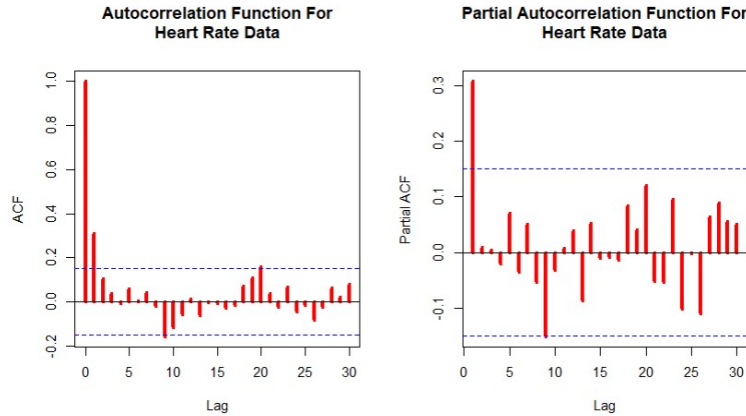


Figure 2.5: ACF And PACF For Heart Rate

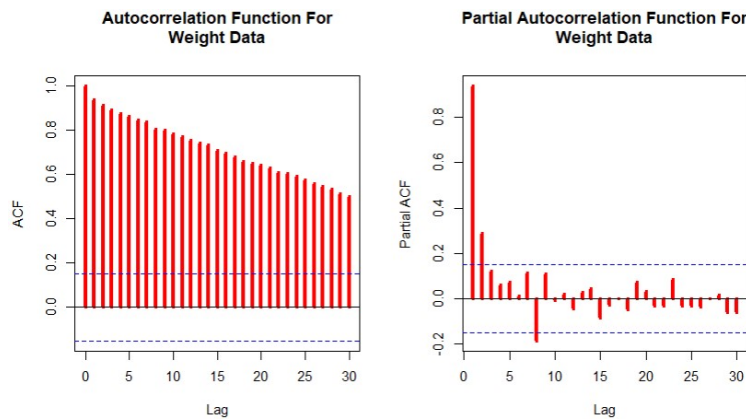


Figure 2.6: ACF And PACF For Weight

As expected, the ACF for the weight data refuses to die out; the data is heavily correlated. After taking a difference, we see the P/ACF plots in Figure 2.7 below.

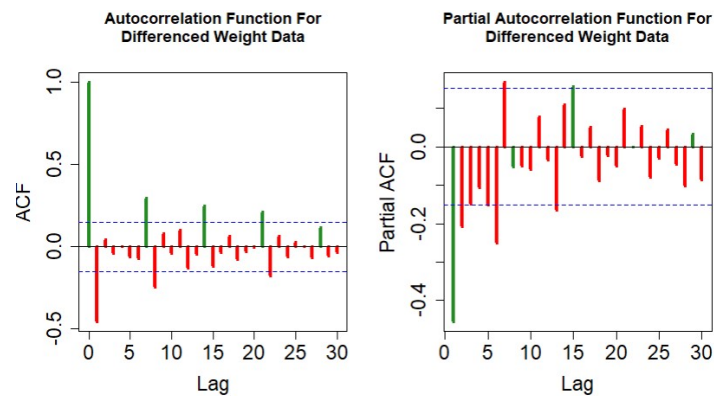


Figure 2.7: ACF And PACF For Differenced Data

2.4 Seasonality

Notice that the ACF for the differenced data in Figure 2.7 above has spikes at lags of 7, 14, and 21 (the seasonal lags are highlighted in green). This indicates that we might try fitting a seasonal component to the data. There is a physical explanation for this seasonality as well— I didn't run on Sunday's and often ate out on the weekend.

To account for this seasonality, we can try to fit a SARIMA with $s = 7$. Since the spikes in the ACF are not growing, we may not need to take a seasonal difference. Nevertheless, we try taking one and see how the P/ACF plots look. They are shown in Figure 2.8 below.

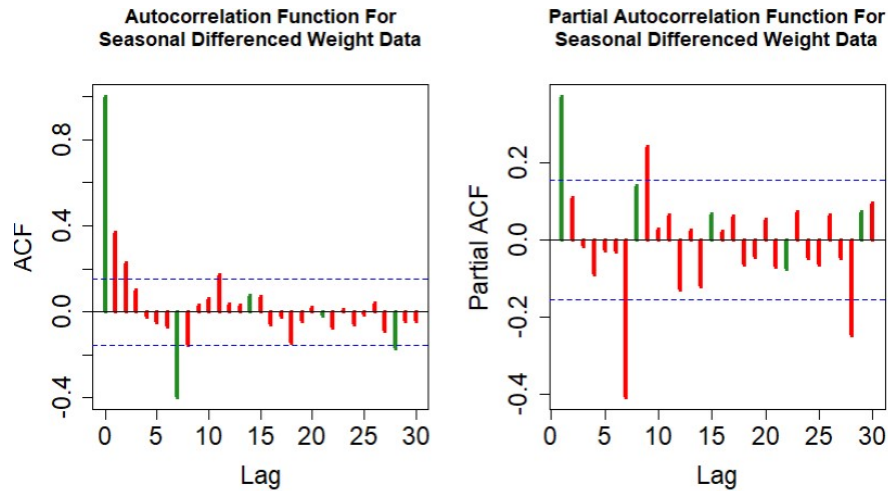


Figure 2.8: ACF And PACF Of Seasonally Differenced Weight Data

There are still some large spikes in the PACF well out into the data, so taking just a seasonal difference may not be sufficient. When taking both a seasonal difference ($D = 1$) and then a regular difference ($d = 1$), we get the plots for our P/ACF in Figure 2.9.

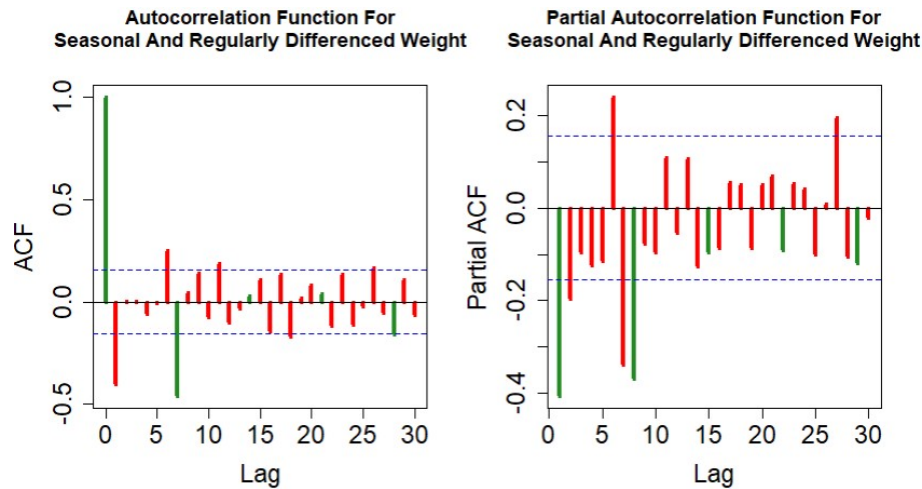


Figure 2.9: ACF And PACF Of Seasonally And Regularly Differenced Weight

2.5 Suggested Models

From Subsection 2.4, there is no clear answer as to which differencing combination (the d and D terms) is best to deal with the weight data.

Based on the ACF and PACF from the regularly differenced data ($d = 1$) from Figure 2.7, we see the seasonal lags of the ACF gradually die out, while the seasonal lags of the PACF have a spike at lag 2 (so, an ARIMA(2,0,0) or ARIMA(2,0,3) for the seasonal component might make sense). The regular lags in the ACF do not completely die out with spikes well out into the series, though the last large spike is at lag 1. The regular lags in the PACF gradually die out, despite having a large spike at lag 6 (so, an ARIMA(0,1,1) for the regular component might make sense). This first bit of our recommended models are ARIMA(0,1,1)(2,0,0)7 and ARIMA(0,1,1)(2,0,3)7.

Based on the ACF and PACF from the seasonally differenced data ($D = 1$) from Figure 2.8, we see the seasonal lags in the ACF die out after a large first spike, while the seasonal lags of the PACF immediately dissipate (so, an ARIMA(0,1,1) or ARIMA(0,1,0) for the seasonal component might make sense). The regular lags in the ACF have meaningful spikes at lags 1 and 2 before having a sinusoidal decay. The regular lags in the PACF do not really die off (so, an ARIMA(0,0,1) or ARIMA(0,0,2) or ARIMA(0,0,3) for the regular component might make sense). The second bit of our recommended models are ARIMA(0,0,1)(0,1,1)7, ARIMA(0,0,1)(0,1,0)7, ARIMA(0,0,2)(0,1,1)7, and ARIMA(0,0,2)(0,1,0)7.

Based on the ACF and PACF from the data that is both seasonally differenced and regularly differenced ($d = 1, D = 1$), we see one large spike in the seasonal lag of the ACF and one large spike in the seasonal lag of the PACF (so, an ARIMA(1,1,0) or ARIMA(0,1,1) or ARIMA(1,1,1) for the seasonal component might make sense). The regular lags in the ACF and PACF do not really die off, but maybe an ARIMA(1,1,0) or ARIMA(0,1,1) or ARIMA(1,1,1) could work. The third and final bit of our recommended models are ARIMA(1,1,1)(1,1,0)7, ARIMA(1,1,1)(0,1,1)7, ARIMA(1,1,1)(1,1,1)7, ARIMA(0,1,1)(1,1,0)7, ARIMA(0,1,1)(0,1,1)7, ARIMA(0,1,1)(1,1,1)7, ARIMA(1,1,0)(1,1,0)7, ARIMA(1,1,0)(0,1,1)7, and ARIMA(1,1,0)(1,1,1)7.

The heart rate data is much simpler. Based on the ACF and PACF for the Heart Rate Data in Figure 2.5, we suggest an AR(1). This is because the last large spike in the ACF and PACF are both at lag one, and the ACF does not completely die off.

2.6 Model Diagnostics

We prioritize the models we identified in subsection 2.5, but we also have lots of computational power to try many different models.

We utilize this power by trying all seasonal ARIMA models with $p, q, P,$ and Q terms less than 5 and d and D terms less than 2. For each of the $5^4 \times 2 \times 2 = 2500$ models, we compute the AIC and BIC for model evaluation, and the p-value from the Ljung-Box Q test to see if the residuals from our model are actually white noise. The top 15 models in terms of BIC are shown in Figure 2.10 below.

```
> df_weight[1:15,]
      ARIMAS_model    aic    bic LBtest
307  ARIMA(0,1,1)(0,1,1)7 509.46 518.71  0.73
807  ARIMA(1,1,1)(0,1,1)7 509.92 522.25  0.81
357  ARIMA(0,1,2)(0,1,1)7 510.49 522.82  0.79
317  ARIMA(0,1,1)(1,1,1)7 511.35 523.68  0.66
308  ARIMA(0,1,1)(0,1,2)7 511.39 523.72  0.69
1307 ARIMA(2,1,1)(0,1,1)7 509.31 524.71  0.88
407  ARIMA(0,1,3)(0,1,1)7 509.43 524.84  0.85
857  ARIMA(1,1,2)(0,1,1)7 510.31 525.72  0.86
327  ARIMA(0,1,1)(2,1,1)7 510.54 525.95  0.67
557  ARIMA(1,0,1)(0,1,1)7 513.60 525.95  0.71
309  ARIMA(0,1,1)(0,1,3)7 510.74 526.14  0.65
817  ARIMA(1,1,1)(1,1,1)7 511.52 526.93  0.74
808  ARIMA(1,1,1)(0,1,2)7 511.66 527.07  0.77
367  ARIMA(0,1,2)(1,1,1)7 512.24 527.65  0.72
358  ARIMA(0,1,2)(0,1,2)7 512.33 527.73  0.75
```

Figure 2.10: ARIMA Model Diagnostics

See that our recommended ARIMA(0,1,1)(0,1,1) had the best BIC. Also notice that all the top models had both a seasonal and regular difference. We finally note that in terms of AIC, the ARIMA(0,1,1)(0,1,1) was also a top performer, and only models with much more terms (e.g. ARIMA(2,1,1)(1,1,4)7), bested it by that metric. Since we are using the models for prediction, we prefer parsimony and so base our decision on the metric that is less forgiving to added parameters; our model choice is the ARIMA(0,1,1)(0,1,1)7.

The heart rate data was stationary to begin with; we only need to consider ARMA models. The top model in terms of BIC (of all combinations of p and q less than 5) was our suggested AR(1). The top ten models are shown in Figure 2.11 below.

```
> df_hr[1:10,]
      ARIMAS_model    aic    bic LBtest
7      ARMA(1,0)  993.33 1002.72  0.86
2      ARMA(0,1)  994.81 1004.20  0.78
8      ARMA(1,1)  995.32 1007.84  0.86
13     ARMA(2,0)  995.32 1007.84  0.86
3      ARMA(0,2)  995.58 1008.10  0.85
4      ARMA(0,3)  997.12 1012.76  0.87
9      ARMA(1,2)  997.32 1012.97  0.86
19     ARMA(3,0)  997.32 1012.97  0.86
14     ARMA(2,1)  997.33 1012.98  0.86
1      ARMA(0,0) 1008.24 1014.50  0.02
```

Figure 2.11: ARIMA Model Diagnostics

3 Parameter Selection

3.1 Heart Rate Data

We fit both series with two models each using the `forecast` package from R. In choosing the coefficients, we are selecting those parameter values which minimize the conditional least squares.

For the heart rate data, the top performing model was our suggested AR(1). Our model is (where $a_t \sim N(0, 20.16)$):

$$a_t = \pi(B)\tilde{Z}_t \quad (3.1)$$

$$a_t \approx (1 - 0.3086B)(Z_t - 55.3846) \quad (3.2)$$

$$Z_t \approx 55.3846 + 0.3086(Z_{t-1} - 55.3846) + a_t \quad (3.3)$$

```
> hr_ar1

Call:
arima(x = ts_hr, order = c(1, 0, 0), method = "ML")

Coefficients:
      ar1  intercept
      0.3086    57.3851
s.e.    0.0731     0.4983

sigma^2 estimated as 20.16:  log likelihood = -493.67,  aic = 993.33
```

Figure 3.1: R Code For Heart Rate Model Coefficients

The second best model was an MA(1). The model fit is (where $a_t \sim N(0, 20.34)$):

$$\tilde{Z}_t = \psi(B)a_t \quad (3.4)$$

$$(Z_t - 55.3846) \approx (1 + 0.2283B)a_t \quad (3.5)$$

$$Z_t \approx 55.3846 + a_t + 0.2283a_{t-1} \quad (3.6)$$

```
> hr_ma1

Call:
arima(x = ts_hr, order = c(0, 0, 1), method = "ML")

Coefficients:
      ma1  intercept
      0.2883    57.3906
s.e.    0.0704     0.4463

sigma^2 estimated as 20.34:  log likelihood = -494.4,  aic = 994.81
```

Figure 3.2: R Code For Heart Rate Model Coefficients

3.2 Weight Data

We give the coefficients to the Seasonal ARIMA models that we fit for the weight data below. The best model in terms of BIC was our suggested SARIMA(0,1,1)(0,1,1). Our model is (where $a_t \sim N(0, 1.268)$):

$$\pi(B)\Pi(B^s)(1-B)^d(1-B^s)^D Z_t = \psi(B)\Psi(B^s)a_t \quad (3.7)$$

$$(1-B)(1-B^7)Z_t \approx (1-0.6659B)(1-0.8147B^7)a_t \quad (3.8)$$

$$Z_t \approx Z_{t-1} + Z_{t-7} - Z_{t-8} + a_t - 0.6659a_{t-1} - 0.8147a_{t-7} + 0.5425a_{t-8} \quad (3.9)$$

```
> weight_sarima011011

call:
arima(x = ts_weight, order = c(0, 1, 1), seasonal = list(order = c(0, 1, 1),
  period = 7), method = "ML")

Coefficients:
      ma1      sma1
    -0.6659   -0.8147
s.e.    0.0826    0.0886

sigma^2 estimated as 1.268:  log likelihood = -251.73,  aic = 509.46
```

Figure 3.3: R Code For Weight Data Model Coefficients

The best model in terms of AIC was a SARIMA(2,1,1)(1,1,4). Our model is (where $a_t \sim N(0, 1.085)$):

$$\pi(B)\Pi(B^s)(1-B)^d(1-B^s)^D Z_t = \psi(B)\Psi(B^s)a_t \quad (3.10)$$

$$(1 + 0.2680B + 0.2111B^2)(1 - 0.7620B^7)(1 - B)(1 - B^7)Z_t \approx \quad (3.11)$$

$$(1 - 0.9142B)(1 - 0.0006B^7 - 0.6255B^{14} - 0.1021B^{21} - 0.2686B^{28})a_t \quad (3.12)$$

```
> weight_sarima211114

call:
arima(x = ts_weight, order = c(2, 1, 1), seasonal = list(order = c(1, 1, 4),
  period = 7), method = "ML")

Coefficients:
      ar1      ar2      ma1      sar1      sma1      sma2      sma3      sma4
    0.2680  0.2111  -0.9142  -0.7620  -0.0006  -0.6255  -0.1021  -0.2686
s.e.    0.1071  0.0985  0.0740  0.1412  0.4828  0.4731  0.1948  0.1551

sigma^2 estimated as 1.085:  log likelihood = -245.62,  aic = 509.25
```

Figure 3.4: R Code For Weight Data Model Coefficients

4 Forecasting

With our top models in hand, we try to forecast our series. We forecast our series out a month, and compare it to our test data that we reserved from the outset. The forecast can be done automatically with R using the `forecast()` function.

The results for the weight data are shown in Figure 4.1 below. The red shading refers to the 95% Prediction Interval for the SARIMA(2,1,1)(1,1,4)₇ model while the blue shading refers to the 95% Prediction Interval for the SARIMA(0,1,1)(0,1,1)₇ model.

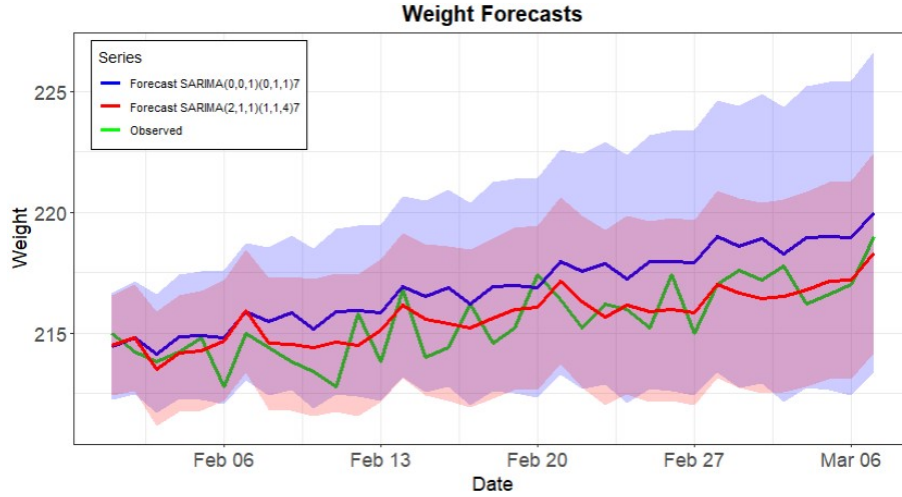


Figure 4.1: Comparison Of Predicted And Observed Values For Weight Data

While an argument could have been made that the more comprehensive model favored by AIC overfits to the training data, it actually does a better job at forecasting our test data compared to the more parsimonious model we suggested (out of sample RMSE of 0.95 compared to out of sample RMSE of 1.73). In either case, see how the prediction bounds grow the larger we move from observed data. This is only natural—our uncertainty about the future grows based on the time.

Since the recommended model for the heart rate data was an AR(1), the forecast will be mean-reverting; the second term in our model $Z_t \approx 55.3846 + 0.3086(Z_{t-1} - 55.3846) + a_t$ becomes smaller and smaller. Our forecasts are:

$$\begin{aligned}\widehat{Z}_{169}(1) &= 55.38462 + 0.3086(59 - 55.38462) \approx 56.500 \\ \widehat{Z}_{169}(2) &= 55.38462 + 0.3086(56.500 - 55.38462) \approx 55.729 \\ \widehat{Z}_{169}(3) &= 55.38462 + 0.3086(55.729 - 55.38462) \approx 55.491 \\ &\vdots\end{aligned}$$

Indeed, after twelve units, our predictions stabilize up to the fifth decimal.

```
> forecast_hr_ar1
```

	date	forecast	lower	upper	observed	resid
1	2023-02-01	56.50081	47.70026	65.30136	58	1.4991863
2	2023-02-02	55.72947	46.51929	64.93965	55	-0.7294680
3	2023-02-03	55.49140	46.24315	64.73965	55	-0.4914009
4	2023-02-04	55.41792	46.16605	64.66979	50	-5.4179241
5	2023-02-05	55.39525	46.14303	64.64746	53	-2.3952464
6	2023-02-06	55.38825	46.13600	64.64049	54	-1.3882471
7	2023-02-07	55.38609	46.13384	64.63834	54	-1.3860869
8	2023-02-08	55.38542	46.13317	64.63767	59	3.6145799
9	2023-02-09	55.38521	46.13296	64.63746	59	3.6147856
10	2023-02-10	55.38515	46.13290	64.63740	61	5.6148491
11	2023-02-11	55.38513	46.13288	64.63738	55	-0.3851312
12	2023-02-12	55.38513	46.13287	64.63738	55	-0.3851252

5 Appendix

```

1- #####Code written By Liam Flaherty For ST534 Final Project#####
2- #####1. Initial Data#####
3- ###1a. Load in data and required packages###
4- library(tidyverse)
5- library(scales)
6- library(forecast)
7-
8- path="C:/Users/LiamFlaherty/Documents/Academics/ST534 Time Series/Project/weight.csv"
9- weight=read.csv(path)
10-
11- weight=weight|>
12-   mutate(Date=as.Date(Date)) |>
13-   mutate(Day=weekdays(Date)) |>
14-   select(Date, Day, weight, Lower, Upper, HR)
15-
16- str(weight)
17- summary(weight)
18-
19-
20-
21- ###1b. Split into training and test###
22- train=weight[which(weight$Date<"2023-02-01"),]
23- test=weight[which(weight$Date>="2023-02-01"),]
24-
25-
26-
27- #####2. Exploratory Data Analysis#####
28- ###2a. Heart Rate###
29- ggplot(train, aes(x=Date, y=HR)) +
30-   geom_line(color="red", linewidth=1) +
31-   labs(title="Evolution Of Heart Rate",
32-        x="Date",
33-        y="Heart Rate") +
34-   scale_x_date(
35-     date_breaks="1 month",
36-     date_labels="%b %Y") +
37-   theme_bw() +
38-   theme(
39-     plot.title=element_text(hjust=0.5, size=16), #Center the title#
40-     axis.text=element_text(size=14),
41-     axis.title=element_text(size=14),
42-     axis.text.x=element_text(angle=45, hjust=1))
43-
44- hist(train$HR,
45-      main="Histogram of HR (8/2022 - 3/2023)",
46-      xlab="Heart Rate",
47-      ylab="Frequency",
48-      xlim=c(45,70),
49-      col="Red")
50-
51- sd(train$HR)
52-
53-
54-
55-
56- ###2b. Weight###
57- ggplot(train, aes(x=Date, y=weight)) +
58-   geom_line(color="blue", linewidth=1) +
59-   labs(title="Evolution of weight",
60-        x="Date",
61-        y="weight") +
62-   scale_x_date(
63-     date_breaks="1 month",
64-     date_labels="%b %Y") +
65-   theme_bw() +
66-   theme(
67-     plot.title=element_text(hjust=0.5, size=16), #Center the title#
68-     axis.text=element_text(size=14),
69-     axis.title=element_text(size=14),
70-     axis.text.x=element_text(angle=45, hjust=1))
71-
72-
73-
74-
75-
76- #####3. Analysis#####
77- ###3a. Convert to time series###
78- ts_weight=ts(train$weight)          #convert to time series object#
79- ts_hr=ts(train$HR)
80-
81-
82-
83- ###3b. White Noise Test###
84- #Clear that weight is not white noise#
85- whitenoise6=Box.test(ts_hr,          #Do we need to fit model?#
86-   lag=6,
87-   type="Ljung-Box")
88- whitenoise6          #p small \implies yes#
89-
90- whitenoise12=Box.test(ts_hr,          #Do we need to fit model?#
91-   lag=12,
92-   type="Ljung-Box")
93- whitenoise12          #p small \implies yes#
94-
95-
96-
97- ###3c. Test for stationarity###
98- adf_result_weight=suppresswarnings(adf.test(ts_weight))
99- adf_result_weight          #difference needed#
100-

```

```

101 adf_result_hr=suppresswarnings(adf.test(ts_hr))
102 adf_result_hr                                     #stationary#
103
104
105
106 ###3c. Initial P/ACF For HR###
107 par(mfrow=c(1,2))                                #split the display to show two figures in one plot#
108
109 acf(ts_hr,
110     main=paste0("Autocorrelation Function For", "\n", "Heart Rate Data"),
111     lag.max=30,
112     ci.col="blue",
113     col="red",
114     lwd=4)
115
116 pacf(ts_hr,
117     main=paste0("Partial Autocorrelation Function For", "\n", "Heart Rate Data"),
118     lag.max=30,
119     ci.col="blue",
120     col="red",
121     lwd=4)
122
123 par(mfrow=c(1,1))                                #back to one figure per plot#
124
125
126
127 ###3d. Initial P/ACF For Weight###
128 par(mfrow=c(1,2))
129
130 acf(ts_weight,
131     main=paste0("Autocorrelation Function For", "\n", "Weight Data"),
132     lag.max=30,
133     ci.col="blue",
134     col="red",
135     lwd=4)
136
137 pacf(ts_weight,
138     main=paste0("Partial Autocorrelation Function For", "\n", "Weight Data"),
139     lag.max=30,
140     ci.col="blue",
141     col="red",
142     lwd=4)
143
144 par(mfrow=c(1,1))
145
146
147
148 ###3e. P/ACF For weight with Regular Difference###
149 train_diff=diff(ts_weight, lag=1)
150
151 par(cex.axis=1.5, cex.lab=1.5)
152 plot(train_diff,
153     main="Time Series Of Differenced Weight Data, Lag=1",
154     xlab="Time",
155     ylab="Difference",
156     cex.axis=2,
157     cex.lab=2,
158     col="blue",
159     lwd=2)
160
161 par(mfrow=c(1,2))
162 acf(train_diff,
163     main=paste0("Autocorrelation Function For", "\n", "Differenced Weight Data"),
164     lag.max=30,
165     ci.col="blue",
166     col=ifelse((0:30 % 7)==0, "forestgreen", "red"),
167     lwd=4)
168
169 pacf(train_diff,
170     main=paste0("Partial Autocorrelation Function For", "\n", "Differenced Weight Data"),
171     lag.max=30,
172     ci.col="blue",
173     col=ifelse((0:30 % 7)==0, "forestgreen", "red"),
174     lwd=4)
175
176 par(mfrow=c(1,1))
177
178
179
180 ###3f. P/ACF For weight with Just Seasonal Difference###
181 ts_weight_sdiff=diff(ts_weight, lag=7)             #just D=1#
182
183 par(mfrow=c(1,2))
184
185 acf(ts_weight_sdiff,
186     main=paste0("Autocorrelation Function For", "\n", "Seasonal Differenced weight Data"),
187     lag.max=30,
188     ci.col="blue",
189     col=ifelse((0:30 % 7)==0, "forestgreen", "red"),
190     lwd=4)
191
192 pacf(ts_weight_sdiff,
193     main=paste0("Partial Autocorrelation Function For", "\n", "Seasonal Differenced weight Data"),
194     lag.max=30,
195     ci.col="blue",
196     col=ifelse((0:30 % 7)==0, "forestgreen", "red"),
197     lwd=4)
198
199 par(mfrow=c(1,1))
200

```



```

201
202
203 ###3g. P/ACF For Weight With Both Differences###
204 ts_weight_both=diff(diff(ts_weight, lag=7), lag=1) #D=1, d=1#
205
206 par(mfrow=c(1,2))
207
208 acf(ts_weight_both,
209     main=paste0("Autocorrelation Function For", "\n", "Seasonal And Regularly Differenced weight"),
210     lag.max=30,
211     ci.col="blue",
212     col=ifelse((0:30 %% 7)==0, "forestgreen", "red"),
213     lwd=4)
214
215 pacf(ts_weight_both,
216     main=paste0("Partial Autocorrelation Function For", "\n", "Seasonal And Regularly Differenced weight"),
217     lag.max=30,
218     ci.col="blue",
219     col=ifelse((0:30 %% 7)==0, "forestgreen", "red"),
220     lwd=4)
221
222 par(mfrow=c(1,1))
223
224
225
226
227
228 #####4. Try A Bunch Of Models#####
229 ###4a. For weight Data###
230 ARIMAS_model_weight=vector()
231 aic_weight=vector()
232 bic_weight=vector()
233 LBtest_weight=vector()
234 s=7 #From Analysis#
235 m=5 #the number of MA and AR terms to try#
236 i=0 #to keep track of iterations
237
238 ~ for (p in 1:m) {
239 ~   for (d in 1:2) {
240 ~     for (q in 1:m) {
241 ~       for (P in 1:m) {
242 ~         for (D in 1:2) {
243 ~           for (Q in 1:m) {
244 ~             i=i+1
245 ~             print(paste0("i=", round(i/(m^4*2^2), 2)) #where we're at in the process#
246 ~
247 ~             mymodel=paste0("ARIMA(", p-1, ", ", d-1, ", ", q-1, ")(", P-1, ", ", D-1, ", ", Q-1, ")"),s)
248 ~
249 ~             setTimeLimit(cpu=3, elapsed=3) #otherwise would take forever#
250 ~             model_result=tryCatch({ #for convergence problems#
251 ~               model=arima(ts_weight,
252 ~                 order=c(p-1,d-1,q-1),
253 ~                 seasonal=list(order=c(P-1,D-1,Q-1), period=s),
254 ~                 method="ML") #By Maximum Likelihood#
255 ~
256 ~               ARIMAS_model_weight[i]=mymodel
257 ~               aic_weight[i]=round(AIC(model),2)
258 ~               bic_weight[i]=round(BIC(model),2)
259 ~               LBtest_weight[i]=round(Box.test(residuals(model), lag=21, type="Ljung-Box")$p.value,2)
260 ~
261 ~             }, error=function(e) { #for convergence problems#
262 ~               ARIMAS_model_weight[i]=mymodel
263 ~               aic_weight[i]=999
264 ~               bic_weight[i]=999
265 ~               LBtest_weight[i]=999
266 ~             })
267 ~             setTimeLimit(cpu=Inf, elapsed=Inf)
268 ~
269 ~           }
270 ~         }
271 ~       }
272 ~     }
273 ~   }
274 ~ }
275
276 df_weight=data.frame(ARIMAS_model=ARIMAS_model_weight,
277                      aic=aic_weight,
278                      bic=bic_weight,
279                      LBtest=LBtest_weight)
280 df_weight=df_weight[which(df_weight$bic<999),]
281 df_weight=df_weight[order(df_weight$bic),]
282 df_weight[1:15,]
283
284 save(df_weight, file="modelfit_weight.R")
285 load("modelfit_weight.R") #so don't have to run this part of the code#
286
287
288
289 ###4b. For HR Data###
290 ARIMAS_model_hr=vector()
291 aic_hr=vector()
292 bic_hr=vector()
293 LBtest_hr=vector()
294 m=5 #the number of MA and AR terms to try#
295 i=0
296
297 ~ for (p in 0:m) {
298 ~   for (q in 0:m) {
299 ~     i=i+1
300 ~     mymodel=paste0("ARMA(", p, ", ", q, ")")

```

```

301
302     model=arima(ts_hr,
303               order=c(p,0,q),
304               method="ML")           #by Maximum Likelihood#
305
306     ARIMAS_model_hr[i]=mymodel
307     aic_hr[i]=round(AIC(model),2)
308     bic_hr[i]=round(BIC(model),2)
309     LBtest_hr[i]=round(Box.test(residuals(model), lag=21, type="Ljung-Box")$p.value,2)
310 }
311 }
312
313 df_hr=data.frame(
314   ARMA_model=ARIMAS_model_hr,
315   aic=aic_hr,
316   bic=bic_hr,
317   LBTest=LBtest_hr)
318 df_hr=df_hr[order(df_hr$bic),]
319 df_hr[1:10,]
320
321 save(df_hr, file="model_fit_hr.R")
322 load("model_fit_hr.R")           #so don't have to run this part of the code#
323
324
325
326 #####4c. Getting parameter weights###
327 hr_ar1=arima(ts_hr,                #best in terms of BIC#
328             order=c(1,0,0),
329             method="ML")
330
331 hr_ma1=arima(ts_hr,                #our recommendation; 2nd best in AIC and BIC#
332             order=c(0,0,1),
333             method="ML")
334
335 weight_sarima211114=arima(ts_weight, #best in terms of AIC#
336                          order=c(2,1,1),
337                          seasonal=list(order=c(1,1,4), period=7),
338                          method="ML")
339
340 weight_sarima011011=arima(ts_weight, #best in terms of BIC; our recommendation#
341                          order=c(0,1,1),
342                          seasonal=list(order=c(0,1,1), period=7),
343                          method="ML")
344
345 hr_ar1
346 hr_ma1
347 weight_sarima211114
348 weight_sarima011011
349
350
351
352 #####5. Forecast#####
353 #####5a. Weight Data###
354 forecast_weight_aic=predict(weight_sarima211114,
355                             n.ahead=nrow(test))
356
357 forecast_weight_aic=data.frame(
358   date=seq(from=test$date[1], to=test$date[nrow(test)], by="day"),
359   forecast=forecast_weight_aic$pred,
360   lower=forecast_weight_aic$pred-1.96*forecast_weight_aic$se,
361   upper=forecast_weight_aic$pred+1.96*forecast_weight_aic$se,
362   observed=test$weight,
363   resid=test$weight-forecast_weight_aic$pred
364 )
365
366 forecast_weight_bic=predict(weight_sarima011011,
367                             n.ahead=nrow(test))
368
369 forecast_weight_bic=data.frame(
370   date=seq(from=test$date[1], to=test$date[nrow(test)], by="day"),
371   forecast=forecast_weight_bic$pred,
372   lower=forecast_weight_bic$pred-1.96*forecast_weight_bic$se,
373   upper=forecast_weight_bic$pred+1.96*forecast_weight_bic$se,
374   observed=test$weight,
375   resid=test$weight-forecast_weight_bic$pred
376 )
377
378 rmse_weight_aic=(sum(forecast_weight_aic$resid^2)/nrow(forecast_weight_aic))^(0.5)
379 rmse_weight_bic=(sum(forecast_weight_bic$resid^2)/nrow(forecast_weight_bic))^(0.5)
380 rmse_weight_aic
381 rmse_weight_bic           #just curious#
382
383 forecast_weight_aic
384 forecast_weight_bic
385
386
387
388 #####5b. Plot Weight Data###
389 ggplot() +
390   geom_line(data=forecast_weight_bic,
391             aes(x=date, y=observed, color="observed"),
392             size=1.2) +
393   geom_line(data=forecast_weight_bic,
394             aes(x=date, y=forecast, color="Forecast SARIMA(0,0,1)(0,1,1)?"),
395             size=1.2) +
396   geom_ribbon(data=forecast_weight_bic,
397              aes(x=date, ymin=lower, ymax=upper),
398              fill="blue",
399              alpha=0.2) +
400   geom_line(data=forecast_weight_aic,

```